US Municipal Records Landscape: Deep Research Report

Part 1: Validating the "30% PDF-Only" Estimate

The Universe of US Local Governments

The US has a staggering number of local governments. The 2017 Census of Governments counted 38,779 general-purpose local governments: 19,495 municipalities, 16,253 townships, and 3,031 counties. On top of that, there are 38,542 special districts and roughly 13,000 school districts — all of which also hold public meetings and produce records.

Of those ~19,500 municipalities, the population distribution is wildly skewed. As of 2019, 16,410 incorporated places (84%) had a population under 10,000. Only about 769 cities exceed 50,000 people, and just 10 have populations over 1 million. This matters enormously for the records question because technology adoption correlates strongly with population and budget.

What We Actually Know About Platform Adoption

Here are the verified client counts from the major vendors:

Platform	Verified Client Count	Source
Granicus (total, all products)	7,000+ government orgs	Granicus website, 2025
Granicus (Legistar specifically)	25 of top 50 US cities; ~1,300 clients for legislative products	Granicus marketing; GovCIO Outlook
CivicPlus (agenda management)	4,200 clerk customers	CivicPlus product page, 2025
CivicPlus (total, all products)	4,000+ local governments	CivicPlus acquisition press, 2021
Municode (codification)	4,250+ local governments	Municode/CivicPlus acquisition press
iCompass	3,000+ local government leaders (LeadOn network)	iCompass website

Critical caveat: These numbers overlap significantly. A city might use Granicus for meeting video, CivicPlus for their website, and Municode for codification. Having a CivicPlus website doesn't mean they use CivicPlus for agenda management. Granicus's 7,000 figure includes all products (communications, web, meetings), not just Legistar.

Attempting a Real Breakdown

Let me work through the math carefully for the ~19,500 municipalities plus ~16,250 townships (35,750 total sub-county general-purpose governments):

Tier 1: Structured API access (Legistar cities)

Legistar has roughly 500-1,300 clients for legislative management. Most of these are cities over 50,000. The Legistar Web API provides structured access to events, matters, persons, votes, and attachments. This is the gold standard for data accessibility.

Estimated coverage: **~800-1,300 governments** (2-4% of total, but covering ~40% of the US population)

Tier 2: Dedicated agenda management platform (scrapable, structured HTML)

CivicPlus claims 4,200 clerk customers for agenda management. Subtracting overlap with Legistar and counting other platforms (iCompass, Civic Clerk, PrimeGov, Novus/iQM2, BoardDocs), we can estimate another 5,000-7,000 governments using some form of dedicated agenda software that produces semi-structured web pages.

Estimated coverage: **~5,000-7,000 governments** (14-20%)

Tier 3: Website with PDFs (no specialized agenda software)

These are governments that have a website — often powered by CivicPlus, Municode, or a generic CMS — where they post meeting minutes and agendas as PDF files linked from a page. They don't use dedicated agenda management software, so documents are uploaded manually. The structure varies wildly between municipalities. Some have well-organized document archives going back decades; others have a single "Meeting Minutes" page with a chronological list of PDFs.

This is likely the largest category. Research on rural digital services confirms that many smaller municipalities provide digital information but without sophisticated platforms. A 2023 study published in *Government Information Quarterly* found that while many rural communities use digital platforms, there are significant "technological and cultural barriers" that limit adoption of more advanced digital services.

Estimated coverage: **~12,000-18,000 governments** (34-50%)

Tier 4: Minimal or no web presence

Small townships and villages with populations under 1,000 (there are thousands of these, particularly in the Midwest and New England). Many lack a functional website entirely, or maintain one with only contact information and meeting schedules. Meeting minutes may only be available on request at the clerk's office. A document scanning company notes that "many smaller municipalities still rely on paper records" and face "limited space, tight budgets, and a growing backlog of paperwork." About 60% of the US population lives in towns under 50,000, yet these areas are "relatively under-studied" in digital government research.

Estimated coverage: **~10,000-15,000 governments** (28-42%)

Revised Estimate

Tier	Description	Est. Governments	% of Total	% of US Pop Covered
1	Structured API (Legistar)	~1,000	~3%	~40%
2	Agenda platform (scrapable HTML)	~6,000	~17%	~25%
3	Website with PDFs	~15,000	~42%	~25%
4	Minimal/no web presence	~13,750	~38%	~10%

The original 30% estimate for "PDF-only" was low. A more defensible estimate is **40-50%** of municipalities post meeting records primarily as PDFs on a basic website, with no dedicated agenda management platform. However, the original estimate isn't far off if you're only counting municipalities that actively post records online at all (excluding the ~38% with minimal web presence).

The good news for your project: By population coverage, the picture is much rosier. Tier 1 + Tier 2 platforms serve roughly 65% of the US population. If you build connectors for Legistar + CivicPlus + a generic PDF scraper, you'd have programmatic access to records for the overwhelming majority of Americans — just not the majority of jurisdictions.

Part 2: Big Local News — What They're Doing and Where You Fit

What Big Local News Is

Big Local News is a program at Stanford University's Journalism and Democracy Initiative, led by Cheryl Phillips (longtime data journalism professor) and co-directed by Serdar Tumgoren (former AP data journalist). Their mission: lower the cost of accountability journalism by gathering hard-to-obtain public data, building open-source tools, and collaborating with newsrooms.

They're not a software company — they're a university-based journalism support organization. Their team is small (about 10-12 people including recent hires) but includes serious talent: Simon Willison (creator of Datasette), Gerald Rich (former AP/Atlantic/Axios), Rosie Cima (2024 Emmy winner), Jake Kara (formerly Washington Post), and Shawn McIntosh (former AJC managing editor) as strategy adviser.

Their Tool Ecosystem

Based on the GitHub org and project documentation, here's their complete tool inventory:

Data Scraping Tools:

Tool	What It Does	Status
civic-scraper	Downloads agendas, minutes, docs from local government sites. Supports CivicPlus, Legistar, Granicus, CivicClerk, PrimeGov	Active (62 stars, 35 open issues, 12 tagged "help wanted"). Updated Feb 2026
court-scraper	Collects case info from US county court sites	Active (76 stars, their most popular repo)
warn-scraper	Downloads WARN Act layoff notices from state websites	Active (37 stars, 53 open issues, 14 "help wanted")
clean-scraper	Harvests police bodycam footage and files (CLEAN project)	Active (10 stars)
local-election-results-etl	Extracts election results from local election officials	Active
usc-crime-reports-scraper	USC Department of Public Safety logs	Active
sync-fed-contracts	Federal government contract cancellations from FPDS	Active
sync-doge-scrape	DOGE claim scraper sync	Active
aphis-inspection-reports	USDA animal inspection data	Active (from data-liberation-project)

Platform Tools:

Tool	What It Does
biglocalnews.org	Data archiving platform for journalists, backed by Stanford Digital Repository
datasette.biglocalnews.org	Datasette instance for exploring hosted data
bln-python-client	Python client for the biglocalnews.org API
upload-files	GitHub Action for uploading to biglocalnews.org
warn-transformer	Consolidates/enriches data from warn-scraper
warn-github-flow	GitHub Action for WARN Act ETL automation
prefect-flow-template	Template for Prefect data pipeline deployment

Analysis Projects:

Project	What It Does
california-wildfires	Wildfire cost investigation
covid-19-tweets	COVID info dissemination by local government
warn-california-analysis	WARN Notices analysis for California
deportation-data	Deportation data project (new, Feb 2026)

Agenda Watch — Their Key Overlap With Your Project

Agenda Watch is their flagship local government tool and the direct overlap with your project. Launched June 2023 in partnership with the Reynolds Journalism Institute (RJI), it's designed to help journalists monitor local government agendas and minutes.

What Agenda Watch does:

- Harvests agendas and minutes daily from agencies across the country
- Lets users search documents by keyword, geography, committee type, and date
- Offers email alerts (real-time, daily, or weekly) when matching documents appear
- Documents are also hosted on DocumentCloud for annotation/publishing
- Integrated with MuckRock's account system (105,000+ registered users, 4,000+ newsrooms)

Current coverage:

- They have a list of 90,000+ government agencies nationwide
- Scrapers written for 2,000+ public agencies (as of late 2023 announcement)
- Started with CivicPlus sites, expanding to Legistar, Granicus, CivicClerk, PrimeGov
- Initial test regions: San Francisco, Atlanta, Philadelphia, Los Angeles

What it uses under the hood:

- `civic-scraper` (their open-source Python library) for the actual data harvesting
- Prefect for pipeline orchestration
- DocumentCloud for document storage and search
- MuckRock for user accounts

As of July 2025, they hired two new engineers (Paige Moody, Jake Kara) specifically to build out Agenda Watch, plus an experienced strategy adviser (Shawn McIntosh) focused on how Agenda Watch can serve local newsrooms. They're clearly scaling this up.

Direct Overlap With Your Project

Your Project	Big Local News / Agenda Watch
Scrape council agendas & minutes	✅ civic-scraper does exactly this
Support CivicPlus, Legistar, Granicus	✅ Already supported (5 platforms)
Parse PDFs for structured data	⚠️ They download PDFs but do NOT parse them for structured vote data, names, etc.
Extract individual votes per official	❌ Not in scope for them — they surface documents, not structured data within them
Track officials across meetings	❌ No official/person database
Cross-reference zoning/development data	❌ Out of scope
Connect to open data portals (Socrata, Opendatasoft)	❌ Not part of civic-scraper
Database of structured political records	❌ They store documents, not extracted structured data
NJ/Hudson County specific coverage	⚠️ Would need to check if NJ agencies are in their list
News article scraping	❌ Not part of their toolkit

Where You Can Contribute to Big Local News

Based on the research, here are the highest-value contribution opportunities:

1. New Platform Connectors for civic-scraper (HIGH VALUE)

civic-scraper currently supports 5 platforms. There are open issues requesting additional platforms, including an issue (#14) for OnBase/Hyland/SIRE. Platforms NOT yet supported that you could contribute:

- **iQM2 / Novus Agenda** — Used by Hoboken and many others. JavaScript-rendered portal, needs Playwright-based scraping. This would directly benefit your Hudson County project AND the broader community.
- **BoardDocs** — Used by many school districts and smaller municipalities.
- **Custom NJ municipal sites** — Jersey City uses a CivicPlus-style CMS but with unique document organization. A NJ-specific connector could cover many Garden State municipalities.

The repo has 35 open issues and 12 specifically tagged "help wanted." This is an actively welcoming project.

2. Structured Data Extraction Layer (HIGHEST VALUE — Major Gap)

This is the single biggest gap in the entire civic data ecosystem, and where your project diverges most from what exists. Nobody — not Big Local News, not Council Data Project, not Open Civic Data — is systematically extracting **structured vote data** from meeting minutes PDFs at scale.

Agenda Watch downloads and indexes the documents. What nobody does is:

- Open the PDF
- Use NLP/LLM to extract "Council member X voted YES on ordinance Y"
- Store that as structured data in a queryable database
- Link it to a person/official entity

If you built an LLM-powered extraction pipeline that takes civic-scraper output (downloaded PDFs) and produces structured vote records, that would be a transformational contribution. You could:

- Build it as a standalone module that consumes civic-scraper assets
- Open-source it with the Big Local News team
- Potentially get it integrated into Agenda Watch

3. NJ Agency Discovery and Site Mapping (MEDIUM VALUE)

civic-scraper maintains a CSV of known CivicPlus sites. Contributing a comprehensive map of NJ municipal websites — which platform each uses, the URL pattern for their agendas/minutes, what date ranges are available — would be immediately useful. Their Agenda Watch platform has a form for requesting new agencies. You could do a systematic survey of all 565 NJ municipalities.

4. Bug Fixes on Existing Scrapers (MEDIUM VALUE, Quick Wins)

Issue #91 shows that the CivicPlus scraper crashes on certain sites due to HTML structure variations. There are similar bugs in the backlog. Fixing these is a practical way to establish yourself as a contributor.

5. Court Scraper NJ Expansion (MEDIUM VALUE)

court-scraper (their most popular repo at 76 stars) scrapes US county court sites. NJ court records might not be covered. Expanding coverage to NJ county courts would dovetail with your local politics interest.

Part 3: Strategic Recommendation

Don't Build a Competing Tool — Build the Missing Layer

The landscape research makes one thing very clear: **document collection is a solved (or nearly solved) problem.** Between civic-scraper, Council Data Project, and the Legistar API, there are working tools to download agendas and minutes from most US municipalities on supported platforms.

What DOESN'T exist is the **intelligence layer**: taking those raw PDFs and HTML documents and extracting structured political data — votes, positions, official actions, developer names, dollar amounts. That's your unique value proposition.

Recommended Architecture (Revised)

```
┌─────────────────────────────────────────┐
│ COLLECTION LAYER (already exists)       │
│ civic-scraper / Legistar API / CDP      │
│ → Downloads raw documents               │
└─────────────────────────────────────────┘
                  │
                  ▼
┌─────────────────────────────────────────┐
│ YOUR CONTRIBUTION: EXTRACTION LAYER     │
│ LLM-powered pipeline that reads PDFs    │
│ and produces structured data:           │
│ • Vote records (who voted how on what)  │
│ • Official actions & positions          │
│ • Developer/applicant mentions          │
│ • Financial amounts & contract details  │
│ • Zoning/variance decisions             │
└─────────────────────────────────────────┘
                  │
                  ▼
```

```
│ STORAGE & QUERY LAYER                    │
│ PostgreSQL + full-text search            │
│ • Canonical schema (OCD-compatible)      │
│ • Per-municipality config (YAML)         │
│ • API for querying across jurisdictions  │
└──────────────────────────────────────────┘
                  │
                  ▼
┌──────────────────────────────────────────┐
│ APPLICATION LAYER                        │
│ • LLM-powered natural language queries   │
│ • Alerts (new votes, new developments)   │
│ • Cross-jurisdictional analysis          │
│ • Export to journalists/researchers      │
└──────────────────────────────────────────┘
```

Concrete Next Steps

1. **Join the Big Local News community.** Email Serdar Tumgoren (zstumgoren on GitHub) or Cheryl Phillips. Explain your project. They're explicitly welcoming contributors and hiring.
2. **Contribute an iQM2 connector to civic-scraper.** This directly serves your Hoboken scraping need AND fills a gap in their platform coverage. It's a tangible, well-scoped contribution.
3. **Build the structured extraction prototype.** Start with Jersey City council minutes PDFs. Use Claude to extract vote records into a structured format. Prove it works on 50 meetings, then generalize.
4. **Map NJ municipal platforms.** Create a comprehensive CSV of all Hudson County (and ideally all NJ) municipalities: what platform they use, URL patterns, available date ranges. Contribute this to civic-scraper's site lists.
5. **Pitch the extraction layer to Big Local News.** Once you have a working prototype, propose it as an Agenda Watch enhancement. They have funding (Knight Foundation, Brown Institute, Lenfest, AP, LA Times) and are actively building this exact kind of capability — they just haven't gotten to the extraction step yet.